

GC Explained

Maciej Paszta

AGENDA

1. Object layout
2. Memory model
3. Generations
4. GC Workflow
5. Different types of GC

Object layout

- SyncBlk or an index to SyncBlk table, stores
 - information about lock - lock(this)
 - index of object in AppDomain
 - hash code
 - COM interop
- TypeHandle - describes the nature of the object's type, reference to MethodTable
 - MethodTable holds information about available methods, default constructor etc. Single instance per every type.
 - MethodDesc - describes a single method and it's IL code.
 - Interface Map / Interface VTable Map - maps type to interface it's implementing
 - static variables
 - EEClass - complimentary to MethodTable; also describes a type, but it contains less frequently used data (usually needed for JIT compilation)
 - Instance fields
 - String literals - actually points to global list of all string l

Memory model

- Each process works in its own virtual address space
 - developer doesn't operate on physical memory
 - can be fragmented - that is the address space contains holes between allocated space
- Stack
 - where value types (int, decimal, enum, struct etc.) are stored.
 - values are removed from the stack as soon as they fall out of scope (stack frame)
 - not subject of Garbage Collection
- GC Heap (or more precisely managed heap)
 - where reference types are stored and allocated.
 - every memory segment allocated on heap represents a type - that's the criteria for determining whether heap is in a valid state
 - heap divided into 2 groups:
 - Small object heap (< 85,000 bytes): Generation 0, Generation 1, Generation 2
 - Large object heap (>= 85,000 bytes)
 - not compacted (in .NET 4.5.1 developer can override this behavior)
 - only collected on full garbage collection runs

Memory model - internals



LoaderHeap (High Frequency Heap)

- CLR artefacts like MethodTable, MethodDesc, FieldDesc - basically the type system.
- static variables - whether they are reference or value types
 - predictable allocation - decrease the chance of fragmentation
 - not a subject of GC



StubHeap

- COM wrapper classes
- P/Invoke wrappers
- Code Access Security

Generations

■ Generation 0

- new objects are allocated here
- most frequently collected generation
- allocated in ephemeral memory segment
- ephemeral generation
- cheap to collect

■ Generation 1

- contains short-lived objects
- acts as a buffer between gen 0 and gen 1
- cheap to collect

■ Generation 2

- long-lived objects
- expensive to collect

Generations

- Flow: Gen0 -> Gen1 -> Gen2
- Ephemeral memory segment:
 - every new memory segment requested by GC
 - designated to store ephemeral generations (Gen0 and Gen1)
 - old ephemeral segments are designated to store Gen2 objects

GC Triggers

- The system is low on physical memory
- The amount of allocated memory on managed heap has surpassed the certain threshold
 - threshold is adjusted in runtime (depending on the survival rate of objects)
- Developer called the `GC.Collect()` method – not advised!
- Developer can prevent GC from running for a small amount of time – `GCSettings.LatencyMode`

GC Phases

1. Mark - marks objects that are still alive, that is they are reachable from GC Roots:


- static variables
- global variables
- stack objects with references to managed heap objects

2. Relocate updates the references to the objects

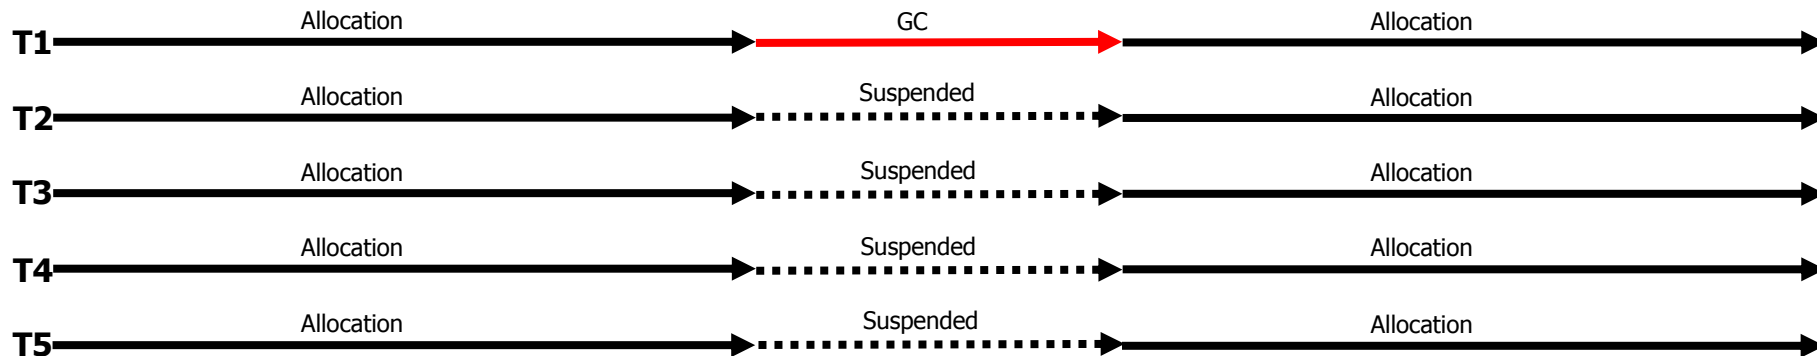
- Objects promoted to higher generation
- Compacted objects

3. Compaction

- reclaims memory occupied by dead objects
- compacts the memory



GC Process



Workstation vs Server GC



Workstation GC:

- Default for desktop and console apps
- Default on 1 CPU machines
- GC runs on user thread (with normal priority), thus must compete with other threads for CPU time
- Collections occur less frequently – improved UI responsiveness
- Can be concurrent (Gen2 collection happens in the background)



Server GC

- Default for services
- Can be turned on by <gcServer> setting
- Dedicated heap and GC thread for each CPU
- GC thread running with highest priority
- Frequent collections – to collect as much temporary objects as possible
- Can be concurrent (Gen2 collection happens in the background)




Concurrent GC swapped for Background workstation/server GC in .NET 4.5.1

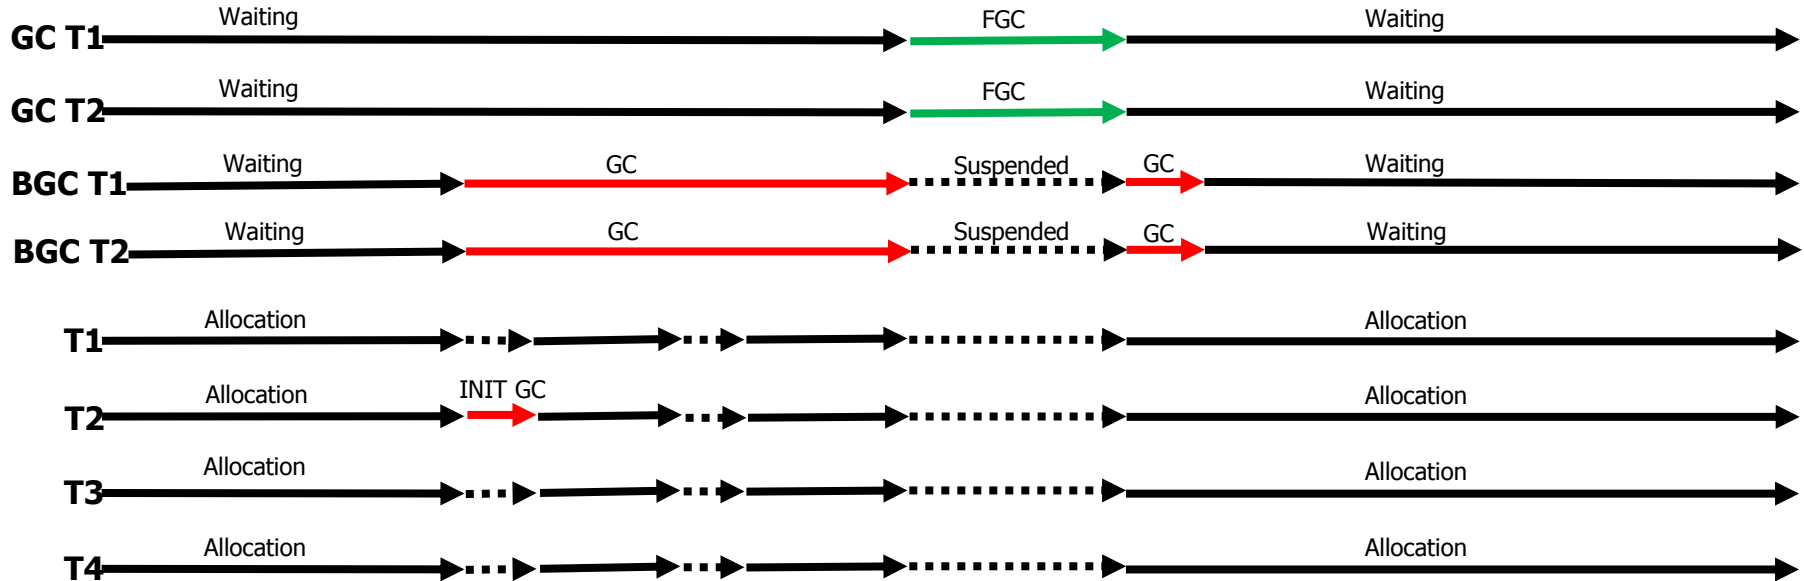
Background GC

- New GC type introduces in .NET 4.5.1
- Dedicated thread to collect Gen2
- Background GC thread doesn't perform compaction
- Gen0 and Gen1 collections require world-stop (they can temporarily suspend BGC collections – this is the main difference when compared to concurrent collection)
- Foreground thread swaps ephemeral segments rather than copying individual objects

Background Workstation GC



Background Server GC



Notes

- GC doesn't affect unmanaged resources

- GC can't move pinned objects

```
byte[] data = new byte[3];  
fixed(byte *ptr = data)  
{  
  
}
```

- Types treated in a special way

- WeakReference
- ConditionalWeakTable<TKey, TValue>

Useful links

- Background and Foreground GC in .NET 4: <http://blogs.msdn.com/b/salvapatuel/archive/2009/06/10/background-and-foreground-gc-in-net-4.aspx>
- GC Fundamentals: [http://msdn.microsoft.com/en-us/library/0xy59wtx\(v=vs.110\).aspx](http://msdn.microsoft.com/en-us/library/0xy59wtx(v=vs.110).aspx)
- Large Object Heap Uncovered: <http://msdn.microsoft.com/en-us/magazine/cc534993.aspx>
- Latency Modes: <https://msdn.microsoft.com/en-us/library/bb384202.aspx>
- Constrained Execution Regions: <https://msdn.microsoft.com/en-us/library/ms228973.aspx>

Thank You!

GC Explained

GFT Poland Sp. z o.o.

Maciej Paszta

Technical Architect

Okraglak, Mielzynskiego 14

61-725 Poznan, Polska

maciej.paszta@gft.com